

Additive Virtual Synthesizer Development in C++

Will Sieber

Department of Computer Science, The College of Wooster

May 27th, 2024

Contents

1	Abstract	3
2	Introduction	4
2.1	Project Goals	4
3	Synthesizer Basics	4
3.1	Components	4
3.1.1	Oscillators	4
3.1.2	Envelope Generators	4
3.2	Extra Components	5
3.3	Additive Synthesis	6
3.4	Extra Synthesis Techniques	6
3.5	Architecture and Prototyping	6
3.5.1	Architecture Diagram	6
4	Implementation	6
4.1	Theory	6
4.1.1	Discrete Signals	6
4.1.2	Buffers	7
4.1.3	Flow of Information	7
4.2	Comparison of Frameworks	7
4.3	VST3 SDK Development	7
4.3.1	Development Tools	7
4.3.2	Workflow and UML diagram	7
4.4	Creating Oscillators	7
4.5	Creating An Envelope Generator	7
5	Conclusion and Future Work	7

1 Abstract

Virtual synthesizers are tools used by musicians to generate sound for use in digital audio workstations. They are repackaged dynamic-link libraries under one of several existing specifications that ensure they are compatible with host applications. Several software development kits and libraries exist with the purpose of developing these tools, but this paper will focus on development under Steinberg's VST3 specification. This will serve as a quick-start guide to those interested in real time audio processing by first providing the theory behind computer audio, and then using this knowledge in a practical setting by developing a virtual additive synthesizer under the VST3 specification. This additive synthesizer will be able to generate several different waveforms with two different oscillators, alongside an envelope generator to control the articulation of the resulting waveform. Additional functionality common to other synthesizers will be included and discussed, but are not required to develop a functional program.

2 Introduction

2.1 Project Goals

3 Synthesizer Basics

3.1 Components

There are several components of a synthesizer that work together in order to generate sound. This section will cover in detail the specific components that were implemented in software. Following this section, there will be another that provides further components that are commonly found in other synthesizers.

3.1.1 Oscillators

Oscillators can be generally broken into two categories: tone generators and controllers. Tone generators repeat periodic waves in order to generate sound. Simple tone generators can be expressed as mathematical expressions, such as $f(x) = \sin(x)$, while complex tone generators allow for any waveform to be repeated, regardless if they can be expressed as a single expression. Complex oscillators generally allow for a number of different waveforms to be represented in a single table, allowing for a user to select any number of waveforms in real time. Controller oscillators, on the other hand, are much lower frequency than tone oscillators, thus have the name *Low-Frequency Oscillators* (LFOs). These will be discussed in brief at a later section. For the purposes of this project, simple oscillators are used as tone generators.

3.1.2 Envelope Generators

In order to provide articulation, envelope generators describe how the amplitude of a particular sound should change over time. Several parameters are used to accomplish this goal, and those are the synthesizer's *attack*, *decay*, *sustain*, and *release*. *Attack* describes how much time it takes for the synthesizer to reach full amplitude upon receiving the signal to generate sound. Note that this is not the s A short attack will reach full amplitude quickly,

such as playing on a piano or organ. Longer attacks will take more time to reach full amplitude, which create a "swelling" effect that is similar to that of a violin increasing in volume. *Decay* describes how much time the synthesizer should take to reach a particular amplitude once the attack is finished. A longer decay will prolong a sound prior to the sustain, and vice versa. *Sustain*, unlike attack and decay, describes the amplitude itself to which the decay will reach upon completion. When reached, the synthesizer will continuously play at this amplitude until a signal is sent that it should stop generating sound. Upon this signal, *release* describes how long the synthesizer should take to stop generating sound by gradually decreasing the amplitude to zero.

In more technical terms, the envelope generator of our synthesizer is its own finite state machine that modifies our signal. The signal is modified by multiplying our source to a constantly changing value. This value has different behavior as to how it changes depending on if it is in the attack state, the decay state, and so on.

One may notice a problem with this model - the rate at which the attack, decay, and release change amplitude itself is not described. Powerful synthesizers allow for this change to be represented as a polynomial function, but for the purposes of this project, the rate at which each parameter changes will simply be linear.

With these two components in mind, one can create a simple synthesizer in the same manner that is described in this project.

3.2 Extra Components

Filters

Digitally Controlled Amplifiers

Low-Frequency Oscillators

LFOs are generally used as an input to automate other parameters of a synthesizer, allowing for a greater range of timbres.

3.3 Additive Synthesis

This section describes the most common forms of synthesis techniques that can be used by musicians to create sound. A complex synthesizer will have functionality for each of the following and then some.

3.4 Extra Synthesis Techniques

Subtractive Synthesis

Wavetable Synthesis

FM Synthesis

3.5 Architecture and Prototyping

3.5.1 Architecture Diagram

4 Implementation

4.1 Theory

There are a number of problems one faces when attempting to process real time audio. They can be summed up by asking the following questions: "How do we create a digital representation of sound?" and "How do we process audio fast enough to ensure there are no unwanted artifacts in our software?". The following section will address these concerns and provide the theory of what is occurring in the background of our software.

4.1.1 Discrete Signals

Sound is represented digitally by taking samples of the original source. An audio sample measures the amplitude of the source wave at a point in time, which is represented in code as a floating point number in the range $[-1, 1]$. Given a piece of audio, the number of samples per second represent the *sample rate* of our recording. As a result, a higher sample rate will result in audio that more closely resembles the original source. Take the following diagram

as an example:

(insert diagram)

Most commonly, audio files will have a sample rate of 44.1 kHz or 48 kHz. That means our software will need to process that number of samples per second, for *each source we have*. For the purposes of this project, there are two audio sources: *OSC1* and *OSC2*. Each oscillator will generate a sound sampled at a rate that is decided by the host application, which by default will be 44.1 kHz.

4.1.2 Buffers

To process this large amount of data quickly, *buffers* are used. A buffer [].

4.1.3 Flow of Information

4.2 Comparison of Frameworks

4.3 VST3 SDK Development

4.3.1 Development Tools

4.3.2 Workflow and UML diagram

4.4 Creating Oscillators

4.5 Creating An Envelope Generator

5 Conclusion and Future Work

References

- [1] M. Jagger. Developing a virtual modular synthesizer for sound waves and midi. *Senior Independent Study Theses*, Paper 9995, 2022.